

```
[root@sandbox lab]# hadoop jar /usr/hdp/2.4.0.0-169/hadoop-mapreduce/hadoop-streaming-2.7.1.2.4.0.0-169.jar -file /root/lab/q1_mapper.py -mapper q1_mapper.py -file /root/lab/q1_reducer.py -reducer q1_reducer.py -input /user/root/dept_salary.txt -output /user/root/q1_streaming_out
WARNING: Use "yarn jar" to launch YARN applications.
16/10/04 15:51:29 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/root/lab/q1_mapper.py, /root/lab/q1_reducer.py] [/usr/hdp/2.4.0.0-169/hadoop-mapreduce/hadoop-streaming-2.7.1.2.4.0.0-169.jar] /tmp/streamjob8106489700785933667.jar tmpDir=null
16/10/04 15:51:34 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/10/04 15:51:35 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050
16/10/04 15:51:36 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/10/04 15:51:36 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050
16/10/04 15:51:39 INFO mapred.FileInputFormat: Total input paths to process : 1
16/10/04 15:51:39 INFO mapreduce.JobSubmitter: number of splits:2
16/10/04 15:51:40 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1475591931336_0001
16/10/04 15:51:42 INFO impl.YarnClientImpl: Submitted application application_1475591931336_0001
16/10/04 15:51:42 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1475591931336_0001/
16/10/04 15:51:42 INFO mapreduce.Job: Running job: job_1475591931336_0001
16/10/04 15:52:12 INFO mapreduce.Job: Job job_1475591931336_0001 running in uber mode : false
16/10/04 15:52:13 INFO mapreduce.Job: map 0% reduce 0%
16/10/04 15:52:34 INFO mapreduce.Job: map 100% reduce 0%
16/10/04 15:52:49 INFO mapreduce.Job: map 100% reduce 100%
16/10/04 15:52:51 INFO mapreduce.Job: Job job_1475591931336_0001 completed successfully
16/10/04 15:52:51 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=15271
    FILE: Number of bytes written=437842
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=20121
    HDFS: Number of bytes written=55
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=2
    Launched reduce tasks=1
    Data-local map tasks=2
    Total time spent by all maps in occupied slots (ms)=37121
    Total time spent by all reduces in occupied slots (ms)=11924
    Total time spent by all map tasks (ms)=37121
    Total time spent by all reduce tasks (ms)=11924
```





Map-Reduce Framework

Map input records=1001
Map output records=1001
Map output bytes=13263
Map output materialized bytes=15277
Input split bytes=226
Combine input records=0
Combine output records=0
Reduce input groups=5
Reduce shuffle bytes=15277
Reduce input records=1001
Reduce output records=5
Spilled Records=2002
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=186
CPU time spent (ms)=2870
Physical memory (bytes) snapshot=517435392
Virtual memory (bytes) snapshot=2528313344
Total committed heap usage (bytes)=354418688

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=19895

File Output Format Counters

Bytes Written=55

16/10/04 15:52:51 INFO streaming.StreamJob: Output directory: /user/root/q1_streaming_out

[root@sandbox lab]# hadoop fs -cat /user/root/q1_streaming_out/part-00000 | tail -n 15

| | |
|-----------|-----|
| Developer | 39 |
| Marketing | 990 |
| QA | 21 |
| Research | 246 |
| Sales | 14 |