

```
[root@sandbox lab]# hadoop jar /usr/hdp/2.4.0.0-169/hadoop-mapreduce/hadoop-streaming-2.7.1.2.4.0.0-169.jar -file /root/lab/q3_mapper.py -mapper q3_mapper.py -file /root/lab/q3_reducer.py -reducer q3_reducer.py -input /user/root/dept_course.txt -output /user/root/l_streaming2_out
WARNING: Use "yarn jar" to launch YARN applications.
16/10/05 02:58:21 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/root/lab/q3_mapper.py, /root/lab/q3_reducer.py] [/usr/hdp/2.4.0.0-169/hadoop-mapreduce/hadoop-streaming-2.7.1.2.4.0.0-169.jar] /tmp/streamjob3547891781051037851.jar tmpDir=null
16/10/05 02:58:24 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/10/05 02:58:25 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050
16/10/05 02:58:26 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/10/05 02:58:26 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050
16/10/05 02:58:27 INFO mapred.FileInputFormat: Total input paths to process : 1
16/10/05 02:58:27 INFO mapreduce.JobSubmitter: number of splits:2
16/10/05 02:58:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1475591931336_0016
16/10/05 02:58:28 INFO impl.YarnClientImpl: Submitted application application_1475591931336_0016
16/10/05 02:58:28 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1475591931336_0016/
16/10/05 02:58:28 INFO mapreduce.Job: Running job: job_1475591931336_0016
16/10/05 02:58:38 INFO mapreduce.Job: Job job_1475591931336_0016 running in uber mode : false
16/10/05 02:58:38 INFO mapreduce.Job: map 0% reduce 0%
16/10/05 02:58:47 INFO mapreduce.Job: map 100% reduce 0%
16/10/05 02:58:55 INFO mapreduce.Job: map 100% reduce 100%
16/10/05 02:58:55 INFO mapreduce.Job: Job job_1475591931336_0016 completed successfully
16/10/05 02:58:55 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=22279
    FILE: Number of bytes written=451858
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=30633
    HDFS: Number of bytes written=8312
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=2
    Launched reduce tasks=1
    Data-local map tasks=2
    Total time spent by all maps in occupied slots (ms)=14549
    Total time spent by all reduces in occupied slots (ms)=5530
    Total time spent by all map tasks (ms)=14549
    Total time spent by all reduce tasks (ms)=5530
    Total vcore-seconds taken by all map tasks=14549
    Total vcore-seconds taken by all reduce tasks=5530
    Total megabyte-seconds taken by all map tasks=3637250
    Total megabyte-seconds taken by all reduce tasks=1382500
  Map-Reduce Framework
    Map input records=1001
    Map output records=1001
    Map output bytes=20271
    Map output materialized bytes=22285
    Input split bytes=226
    Combine input records=0
    Combine output records=0
    Reduce input groups=988
    Reduce shuffle bytes=22285
    Reduce input records=1001
    Reduce output records=467
    Spilled Records=2002
    Shuffled Maps =2
    Failed Shuffles=0
    Merged Map outputs=2
    GC time elapsed (ms)=160
```

Map-Reduce Framework

Map input records=1001
Map output records=1001
Map output bytes=20271
Map output materialized bytes=22285
Input split bytes=226
Combine input records=0
Combine output records=0
Reduce input groups=988
Reduce shuffle bytes=22285
Reduce input records=1001
Reduce output records=467
Spilled Records=2002
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=160
CPU time spent (ms)=2810
Physical memory (bytes) snapshot=548683776
Virtual memory (bytes) snapshot=2528329728
Total committed heap usage (bytes)=357564416

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=30407

File Output Format Counters

Bytes Written=8312

16/10/05 02:58:55 INFO streaming.StreamJob: Output directory: /user/root/l_streaming2_out

[root@sandbox lab]# hadoop fs -cat /user/root/l_streaming2_out/part-00000 | tail -n 15

```
9132 QA Math
9142 QA History
9153 Research Math
9157 Marketing Math
9219 QA Math
9263 Marketing Math
932 Research Math
944 Sales Math
9481 QA Math
9674 Marketing Math
968 Research Math
9755 QA History
9787 Research Math
9965 Marketing Math
9968 QA History
```

[root@sandbox lab]# hadoop fs -cat /user/root/l_streaming2_out/part-00000 | tail -n 15